

Step 1: Importing Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import random
from datetime import datetime, timedelta
```

Step 2: Creating Dataset Creation

This Dataset represent e-commerce transactions from an online retail store. It includes the following fields:

1. CustomerID.
2. OrderID.
3. Product.
4. Quantity.
5. UnitPrice.
6. PurchaseDate.
7. Country.

```
num_records = 1000 # Adjust size as needed

products = ["Laptop", "Smartphone", "Headphones", "Smartwatch",
            "Tablet", "Gaming Console", "Keyboard", "Mouse", "Monitor", "Printer"]

start_date = datetime(2022, 1, 1)
end_date = datetime(2022, 12, 31)

countries = ["USA", "Canada", "UK", "Germany", "France", "Australia",
            "Japan", "India"]

df = pd.DataFrame({
    "CustomerID": np.random.randint(1000, 2000, num_records),
    "OrderID": np.arange(5000, 5000 + num_records),
    "Product": [random.choice(products) for _ in range(num_records)],
    "Quantity": np.random.randint(1, 11, num_records),
    "UnitPrice": np.random.uniform(5, 500, num_records).round(2),
    "PurchaseDate": [start_date + timedelta(days=random.randint(0,
364)) for _ in range(num_records)],
    "Country": [random.choice(countries) for _ in range(num_records)]
})

# random: Python's random module used for generating random choices
# np.random.randint: Generates random integers for CustomerID and
```

```

Quantity
# np.arange: Generates a range of OrderID values
# timedelta: Used to add random days to the start_date for
PurchaseDate
# for _ in range: List comprehension to generate random choices for
Product and Country
# random.choice: Randomly selects an item from products and countries
lists
# random.randint: Generates random integers for days to add to
start_date
# np.random.uniform: Generates random float values for UnitPrice
# .round: Rounds the UnitPrice to 2 decimal places
# .head: Displays the first 5 rows of the DataFrame

```

```
df.head()
```

	CustomerID	OrderID	Product	Quantity	UnitPrice	PurchaseDate
0	1463	5000	Gaming Console	7	141.04	2022-02-12
1	1418	5001	Keyboard	4	110.79	2022-12-12
2	1056	5002	Headphones	5	312.16	2022-10-19
3	1582	5003	Keyboard	5	76.14	2022-02-20
4	1058	5004	Laptop	5	443.32	2022-02-24

	Country
0	Australia
1	UK
2	Canada
3	India
4	India

1. Data Cleaning and Preparation

```

# 1. Check for missing values
print(df.isnull().sum())

# if there are missing values
df.fillna({"Product": "Unknown", "Quantity": 1, "UnitPrice": 0},
inplace=True)

CustomerID    0
OrderID       0
Product       0

```

```
Quantity      0
UnitPrice     0
PurchaseDate   0
Country        0
dtype: int64
```

```
# 2. Check for duplicates
```

```
print(df.duplicated().sum())
```

```
# if there are duplicates
```

```
df.drop_duplicates(inplace=True)
```

```
0
```

```
# Convert PurchaseDate to datetime
```

```
df["PurchaseDate"] = pd.to_datetime(df["PurchaseDate"])
```

```
# Summary of dataset verification
```

```
print(df.info())
```

```
print(df.describe())
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1000 entries, 0 to 999
```

```
Data columns (total 7 columns):
```

#	Column	Non-Null Count	Dtype
0	CustomerID	1000 non-null	int32
1	OrderID	1000 non-null	int64
2	Product	1000 non-null	object
3	Quantity	1000 non-null	int32
4	UnitPrice	1000 non-null	float64
5	PurchaseDate	1000 non-null	datetime64[ns]
6	Country	1000 non-null	object

```
dtypes: datetime64[ns](1), float64(1), int32(2), int64(1), object(2)
```

```
memory usage: 47.0+ KB
```

```
None
```

	CustomerID	OrderID	Quantity	UnitPrice	\
count	1000.000000	1000.000000	1000.000000	1000.000000	
mean	1487.932000	5499.500000	5.543000	255.76674	
min	1001.000000	5000.000000	1.000000	5.39000	
25%	1231.500000	5249.750000	3.000000	132.22750	
50%	1484.500000	5499.500000	6.000000	253.53000	
75%	1743.000000	5749.250000	8.000000	384.84500	
max	1999.000000	5999.000000	10.000000	499.65000	
std	294.477484	288.819436	2.898716	144.55164	

	PurchaseDate
count	1000
mean	2022-07-04 12:47:31.199999744
min	2022-01-01 00:00:00
25%	2022-04-05 00:00:00

50%	2022-07-08 00:00:00
75%	2022-10-03 06:00:00
max	2022-12-31 00:00:00
std	NaN

2. Filtering Data

```
# 2.1 Customer and Order Filters:
# 1. Filter all transactions made by customers from a specific country
(e.g., 'USA').
```

```
print((df[df["Country"] == "USA"]).head(10))
```

	CustomerID	OrderID	Product	Quantity	UnitPrice	PurchaseDate
Country						
8	1304	5008	Headphones	3	384.83	2022-03-04
USA						
17	1286	5017	Mouse	4	487.76	2022-08-06
USA						
18	1634	5018	Printer	2	26.79	2022-01-08
USA						
26	1876	5026	Laptop	10	147.46	2022-07-02
USA						
27	1001	5027	Headphones	3	282.72	2022-09-13
USA						
32	1374	5032	Laptop	10	117.94	2022-01-20
USA						
38	1117	5038	Smartwatch	10	315.70	2022-02-04
USA						
40	1697	5040	Tablet	3	173.91	2022-07-14
USA						
41	1648	5041	Printer	3	143.76	2022-08-08
USA						
57	1253	5057	Keyboard	3	35.53	2022-10-04
USA						

```
# 2. Extract orders where the total spend (Quantity * UnitPrice)
exceeds $500.
```

```
df["Total"] = df["Quantity"] * df["UnitPrice"]
```

```
print(df[df["Total"] > 500].head(5))
```

```
# Note: This could be done using MySQL right??
```

	CustomerID	OrderID	Product	Quantity	UnitPrice	PurchaseDate
0	1463	5000	Gaming Console	7	141.04	2022-02-12
2	1056	5002	Headphones	5	312.16	2022-10-19

4	1058	5004	Laptop	5	443.32	2022-02-24
5	1033	5005	Keyboard	8	245.41	2022-12-03
6	1905	5006	Laptop	8	298.40	2022-03-30

	Country	Total
0	Australia	987.28
2	Canada	1560.80
4	India	2216.60
5	Australia	1963.28
6	India	2387.20

```
# 3. Identify customers who purchased more than 3 different products.
customer_count = df.groupby("CustomerID")["Product"].nunique()
print(customer_count[customer_count > 3].head(10))
```

nunique(): Return series Number of unique values within each group

```
CustomerID
1253    4
1257    6
1373    4
1392    4
1408    5
1514    4
1579    4
1582    5
1937    4
1949    4
Name: Product, dtype: int64
```

2.2 Time-Based Filters:

1. Filter transactions that occurred in July 2022.

```
print(df[(df["PurchaseDate"].dt.month == 7) &
(df["PurchaseDate"].dt.year == 2022)].head())
```

	CustomerID	OrderID	Product	Quantity	UnitPrice	PurchaseDate
10	1475	5010	Printer	9	286.76	2022-07-22
16	1642	5016	Tablet	8	463.10	2022-07-26
20	1958	5020	Gaming Console	4	49.33	2022-07-13
26	1876	5026	Laptop	10	147.46	2022-07-02
40	1697	5040	Tablet	3	173.91	2022-07-14

	Country	Total
10	Australia	2580.84
16	Australia	3704.80
20	Canada	197.32
26	USA	1474.60
40	USA	521.73

2. Extract orders placed during weekends.

```
print(df[df["PurchaseDate"].dt.dayofweek.isin([5,6])].head())
```

isin(): Return a boolean Series showing whether each element in the Series is exactly contained in the passed sequence of values.

dt.dayofweek: The day of the week with Monday=0, Sunday=6

dt: Accessor object for datetime like properties of the Series values.

	CustomerID	OrderID	Product	Quantity	UnitPrice	PurchaseDate \
0	1463	5000	Gaming Console	7	141.04	2022-02-12
3	1582	5003	Keyboard	5	76.14	2022-02-20
5	1033	5005	Keyboard	8	245.41	2022-12-03
11	1022	5011	Monitor	8	136.17	2022-01-09
17	1286	5017	Mouse	4	487.76	2022-08-06

	Country	Total
0	Australia	987.28
3	India	380.70
5	Australia	1963.28
11	Australia	1089.36
17	USA	1951.04

3. Identify transactions during specific sales events, like Black Friday or Cyber Monday.

```
black_friday = df[(df["PurchaseDate"].dt.month == 11) &
(df["PurchaseDate"].dt.weekday == 4) &
(df["PurchaseDate"].dt.day >= 23) &
(df["PurchaseDate"].dt.day <= 29)]
print(black_friday)
```

	CustomerID	OrderID	Product	Quantity	UnitPrice	PurchaseDate
167	1892	5167	Keyboard	4	397.68	2022-11-25
729	1374	5729	Smartphone	2	335.67	2022-11-25

	Country	Total
167	India	1590.72
729	France	671.34

3. Sorting Data

```
# 1. Sort transactions by:
# - Total spend in descending order.
print(df.sort_values("Total", ascending=False).head())
```

	CustomerID	OrderID	Product	Quantity	UnitPrice	PurchaseDate \
877	1017	5877	Smartwatch	10	497.35	2022-12-01
779	1455	5779	Keyboard	10	494.19	2022-10-11
711	1548	5711	Monitor	10	491.39	2022-03-30
207	1583	5207	Printer	10	487.46	2022-07-08
484	1656	5484	Gaming Console	10	487.40	2022-03-01

	Country	Total
877	Australia	4973.5
779	France	4941.9
711	Germany	4913.9
207	Australia	4874.6
484	Australia	4874.0

```
# - Purchase date in ascending order.
print(df.sort_values(by="PurchaseDate", ascending=True).head())
```

```
# sort_values(): Sort by the values along either axis.
```

	CustomerID	OrderID	Product	Quantity	UnitPrice	PurchaseDate
715	1462	5715	Smartphone	9	288.79	2022-01-01
871	1447	5871	Printer	6	123.01	2022-01-01
872	1153	5872	Laptop	2	113.52	2022-01-01
557	1304	5557	Keyboard	6	444.09	2022-01-02
499	1053	5499	Smartwatch	3	420.69	2022-01-03

	Country	Total
715	USA	2599.11
871	UK	738.06
872	Canada	227.04
557	USA	2664.54
499	Australia	1262.07

```
# - Product name alphabetically.
print(df.sort_values("Product").head())
```

	CustomerID	OrderID	Product	Quantity	UnitPrice	PurchaseDate \
0	1463	5000	Gaming Console	7	141.04	2022-02-12
236	1389	5236	Gaming Console	8	202.92	2022-04-29
247	1074	5247	Gaming Console	10	197.79	2022-06-21
254	1026	5254	Gaming Console	3	124.08	2022-10-11
804	1318	5804	Gaming Console	1	246.47	2022-01-16

	Country	Total
0	Australia	987.28
236	Canada	1623.36
247	USA	1977.90
254	Germany	372.24
804	Germany	246.47

Will Continue in MySQL

```
df.to_csv("ecommerce_data.csv", index=False)
```

Going Back to Visualization

```
data = pd.read_csv('ecommerce_data_from_SQL.csv', delimiter=';')
data
```

	CustomerID	OrderID	Product	Quantity	UnitPrice	\
0	1197	5000	Mouse	2	331.04	
1	1599	5001	Smartwatch	7	237.26	
2	1584	5002	Smartphone	5	36.82	
3	1779	5003	Monitor	4	13.17	
4	1660	5004	Monitor	8	224.16	
..	

995	1907	5995	Smartwatch	7	215.17
996	1081	5996	Laptop	2	45.93
997	1903	5997	Laptop	3	300.15
998	1506	5998	Printer	6	268.08
999	1582	5999	Keyboard	7	471.05

	PurchaseDate	Country	TotalSpend
0	2022-10-18 00:00:00	India	662.08
1	2022-11-07 00:00:00	Germany	1660.82
2	2022-02-20 00:00:00	Japan	184.10
3	2022-03-19 00:00:00	Japan	52.68
4	2022-04-11 00:00:00	France	1793.28
...
995	2022-11-28 00:00:00	Japan	1506.19
996	2022-05-27 00:00:00	Japan	91.86
997	2022-11-08 00:00:00	Japan	900.45
998	2022-02-14 00:00:00	USA	1608.48
999	2022-11-24 00:00:00	USA	3297.35


[1000 rows x 8 columns]

6. Visualization(using matplotlib) :Extra grad

```
# 1. Plot the total revenue per country using a bar chart.
revenue_country = data.groupby("Country")["TotalSpend"].sum()
# Plot the total revenue per country using a bar chart
revenue_country.plot(kind="bar", color="Blue")

# Add title and labels to the plot
plt.title("Total Revenue per Country")
plt.ylabel("Total Revenue")
plt.xlabel("Country")

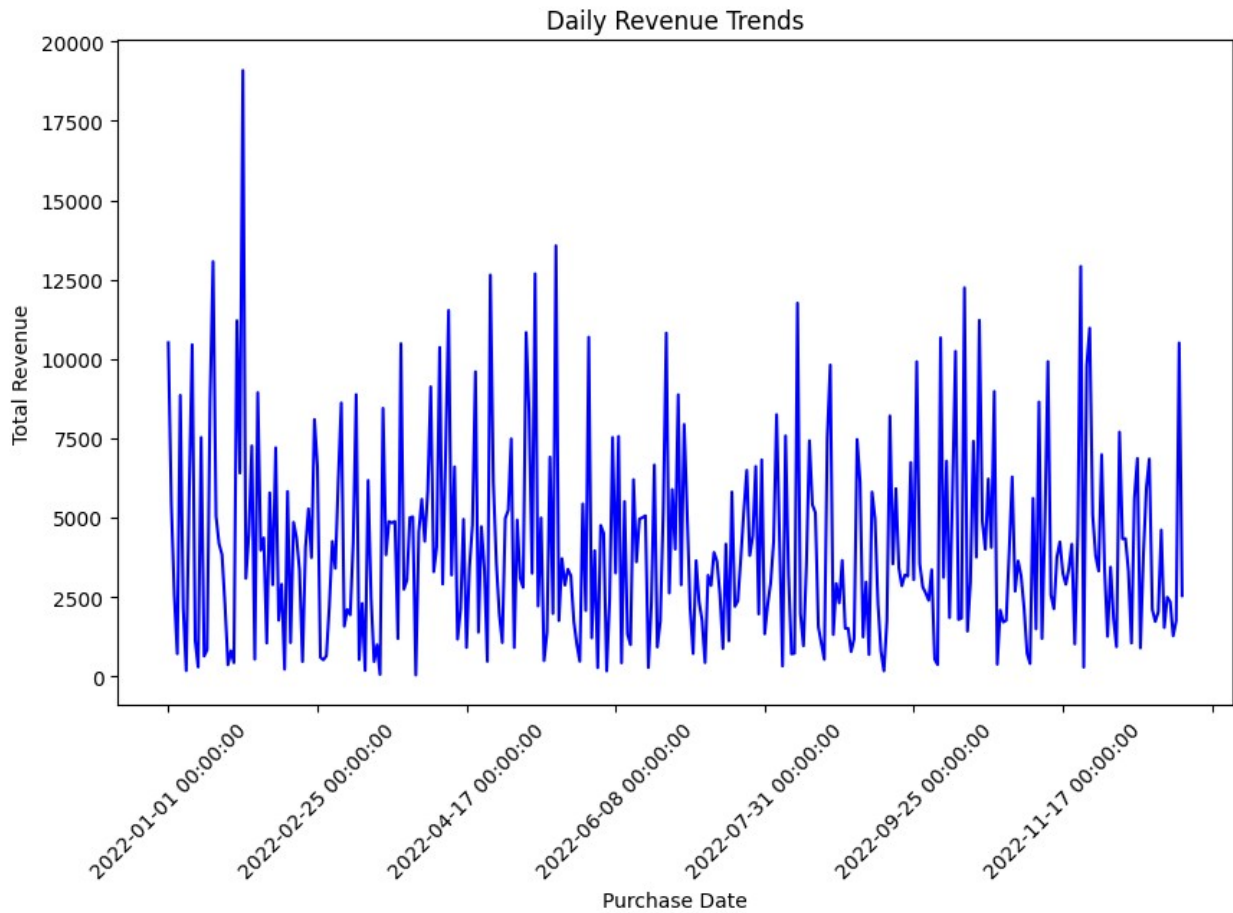
# Display the plot
plt.show()
```



```
# 2. Visualize daily revenue trends with a line chart.
daily_revenue = data.groupby("PurchaseDate")["TotalSpend"].sum()
plt.figure(figsize=(10, 6)) # Set the figure size
daily_revenue.plot(kind="line", color="blue")
plt.title("Daily Revenue Trends")
plt.ylabel("Total Revenue")
plt.xlabel("Purchase Date")

plt.xticks(rotation=45)

plt.show()
```



```
# 3. Use a pie chart to show the distribution of products sold.  
Product_distribution = data["Product"].value_counts()  
plt.pie(Product_distribution, labels = Product_distribution.index,  
startangle = 90)  
plt.show()
```

